

Boston_Real_Estate

September 19, 2022

[Link a repositorio en GitHub](#)

1 Importar bibliotecas

```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

1.1 Importar módulos de Scikit-learn

```
[ ]: from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, GridSearchCV

from sklearn.neural_network import MLPRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, AdaBoostRegressor,
    GradientBoostingRegressor, HistGradientBoostingRegressor
from sklearn.svm import SVR, NuSVR
from sklearn.gaussian_process import GaussianProcessRegressor
```

2 Importar Dataset

- **CRIM**: per capita crime rate by town
- **ZN**: proportion of residential land zoned for lots over 25,000 sq.ft.
- **INDUS**: proportion of non-retail business acres per town
- **CHAS**: Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- **NOX**: nitric oxides concentration (parts per 10 million)
- **RM**: average number of rooms per dwelling
- **AGE**: proportion of owner-occupied units built prior to 1940
- **DIS**: weighted distances to five Boston employment centres
- **RAD**: index of accessibility to radial highways
- **TAX**: full-value property-tax rate per \$10,000
- **PTRATIO**: pupil-teacher ratio by town
- **B**: $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town
- **LSTAT**: % lower status of the population

- **MEDV**: Median value of owner-occupied homes in \$1000's

```
[ ]: url = "https://raw.githubusercontent.com/crisb-7/BostonRealEstate/main/
      ↪bostonRealEstate.csv"
```

```
[ ]: df = pd.read_csv(url)
```

```
[ ]: df.head()
```

```
[ ]:
      CRIM      ZN  INDUS  CHAS    NOX     RM   AGE     DIS  RAD  TAX  PTRATIO  \
0  0.00632  18.0    2.31     0  0.538  6.575  65.2  4.0900    1  296     15.3
1  0.02731   0.0    7.07     0  0.469  6.421  78.9  4.9671    2  242     17.8
2  0.02729   0.0    7.07     0  0.469  7.185  61.1  4.9671    2  242     17.8
3  0.03237   0.0    2.18     0  0.458  6.998  45.8  6.0622    3  222     18.7
4  0.06905   0.0    2.18     0  0.458  7.147  54.2  6.0622    3  222     18.7

      B  LSTAT  MEDV
0  396.90   4.98  24.0
1  396.90   9.14  21.6
2  392.83   4.03  34.7
3  394.63   2.94  33.4
4  396.90   5.33  36.2
```

3 Exploración del dataset

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 511 entries, 0 to 510
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
0   CRIM        511 non-null    float64
1   ZN          511 non-null    float64
2   INDUS       511 non-null    float64
3   CHAS        511 non-null    int64
4   NOX         511 non-null    float64
5   RM          506 non-null    float64
6   AGE         511 non-null    float64
7   DIS         511 non-null    float64
8   RAD         511 non-null    int64
9   TAX         511 non-null    int64
10  PTRATIO     511 non-null    float64
11  B           511 non-null    float64
12  LSTAT       511 non-null    float64
13  MEDV        511 non-null    float64
dtypes: float64(11), int64(3)
```

memory usage: 56.0 KB

Al ver que se tienen solo 5 registros con valores nulos para RM, se quitan estas filas para tener un conjunto de datos homogéneo.

```
[ ]: df = df.dropna(axis = 0)
df.shape
```

```
[ ]: (506, 14)
```

```
[ ]: df.describe()
```


```
[ ]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.617404	11.289526	11.174842	0.069170	0.555209	6.287589
std	8.600123	23.325350	6.824592	0.253994	0.115611	0.703802
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000
25%	0.082268	0.000000	5.190000	0.000000	0.449000	5.885500
50%	0.266005	0.000000	9.690000	0.000000	0.538000	6.209000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.629750
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000

	AGE	DIS	RAD	TAX	PTRATIO	B \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	68.555731	3.775231	9.531621	408.330040	18.498419	356.228379
std	28.161573	2.096147	8.716661	168.382685	2.202078	91.253462
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000
25%	45.025000	2.098500	4.000000	280.250000	17.400000	374.687500
50%	77.500000	3.122200	5.000000	330.000000	19.100000	391.260000
75%	93.975000	5.117675	24.000000	666.000000	20.200000	396.210000
max	100.000000	12.126500	24.000000	711.000000	23.000000	396.900000

	LSTAT	MEDV
count	506.000000	506.000000
mean	12.872569	22.711858
std	7.823528	9.520520
min	1.730000	5.000000
25%	6.950000	17.025000
50%	11.465000	21.200000
75%	17.107500	25.075000
max	76.000000	67.000000

```
[ ]: sns.set(rc={'figure.figsize':(16, 9)})
plt.rcParams["figure.dpi"] = 150
sns.heatmap(df.corr(), annot = True)
plt.show()
```



```
[ ]: # sns.pairplot(df)
      # plt.show()
```

4 Preprocesamiento de los datos

```
[ ]: scaler = StandardScaler()
```

```
[ ]: x = scaler.fit_transform(df.drop(columns = "MEDV"))
      y = df.MEDV
```

5 División train-test

```
[ ]: x_train, x_test, y_train, y_test = train_test_split(x, y, train_size = 0.90,
      random_state = 0)
```

6 Model Cross-validation

```
[ ]: randomState = 0

      Regressors = []
```

```

# Regressors.append(MLPRegressor(random_state = randomState, activation = "relu",
    solver = "adam",
#         hidden_layer_sizes = (100,), alpha = 0.0001,
    learning_rate = "constant",
#         learning_rate_init = 0.0005, max_iter = 5000))

Regressors.append(MLPRegressor(random_state = randomState, activation = "relu",
    solver = "adam",
        hidden_layer_sizes = (100,), alpha = 0.0101, learning_rate = "adaptive",
        learning_rate_init = 0.1, max_iter = 1000))

Regressors.append(KNeighborsRegressor(n_neighbors = 2, weights = "uniform", p = 1))

Regressors.append(DecisionTreeRegressor(random_state=randomState))

Regressors.append(RandomForestRegressor(n_estimators = 250, max_depth = 7,
    random_state=randomState))

Regressors.append(SVR(C = 40.7, epsilon=0.56))

Regressors.append(NuSVR(C = 31.2, nu=0.5))

Regressors.append(AdaBoostRegressor(random_state = randomState))

Regressors.append(GradientBoostingRegressor(random_state = randomState))

Regressors.append(GaussianProcessRegressor(random_state = randomState))

Regressors.append(HistGradientBoostingRegressor(random_state = randomState))

cv_results = []
cv_train_score = []
for regressor in Regressors:
    cv_train=regressor.fit(x_train, y = y_train)
    cv_train_score.append(regressor.score(x_train, y_train))
    cv_results.append(regressor.score(x_test, y_test))

cv_res = pd.DataFrame({"Algorithm":["NN", "KN", "Decision Tree", "Random Forest", "SVR", "NuSVR", "AdaBoost", "GradientBoost", "GaussianProcess", "HistGradientBoosting"],
    "TrainScore":cv_train_score, "TestScore":cv_results})
cv_res.sort_values(by = "TestScore", ascending = False)

```

```

[ ]:           Algorithm  TrainScore  TestScore
7      GradientBoost      0.971574   0.909959
3      Random Forest      0.952958   0.854074
0              NN          0.897917   0.823250
9  HistGradientBoosting    0.958757   0.808177
1              KN          0.929634   0.794065
2      Decision Tree       1.000000   0.778517
4              SVR          0.913308   0.746149
5              NuSVR        0.904299   0.719656
6              AdaBoost     0.862871   0.644497
8      GaussianProcess     1.000000  -0.304075

```