

# GradientBoost\_Boston\_Real\_Estate

September 19, 2022

[Link a repositorio en GitHub](#)

## 1 Importar bibliotecas

```
[ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import uniform as sp_randFloat
from scipy.stats import randint as sp_randInt
%matplotlib inline
%pip install mlxtend --upgrade
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>  
Requirement already satisfied: mlxtend in /usr/local/lib/python3.7/dist-packages (0.14.0)  
Collecting mlxtend  
  Downloading mlxtend-0.21.0-py2.py3-none-any.whl (1.3 MB)  
    || 1.3 MB 23.7 MB/s  
Requirement already satisfied: scipy>=1.2.1 in /usr/local/lib/python3.7/dist-packages (from mlxtend) (1.7.3)  
Requirement already satisfied: joblib>=0.13.2 in /usr/local/lib/python3.7/dist-packages (from mlxtend) (1.1.0)  
Requirement already satisfied: pandas>=0.24.2 in /usr/local/lib/python3.7/dist-packages (from mlxtend) (1.3.5)  
Requirement already satisfied: scikit-learn>=1.0.2 in /usr/local/lib/python3.7/dist-packages (from mlxtend) (1.0.2)  
Requirement already satisfied: matplotlib>=3.0.0 in /usr/local/lib/python3.7/dist-packages (from mlxtend) (3.2.2)  
Requirement already satisfied: setuptools in /usr/local/lib/python3.7/dist-packages (from mlxtend) (57.4.0)  
Requirement already satisfied: numpy>=1.16.2 in /usr/local/lib/python3.7/dist-packages (from mlxtend) (1.21.6)  
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=3.0.0->mlxtend) (2.8.2)  
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=3.0.0->mlxtend) (3.0.9)  
Requirement already satisfied: kiwisolver>=1.0.1 in

```

/usr/local/lib/python3.7/dist-packages (from matplotlib>=3.0.0->mlxtend) (1.4.4)
Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.7/dist-
packages (from matplotlib>=3.0.0->mlxtend) (0.11.0)
Requirement already satisfied: typing-extensions in
/usr/local/lib/python3.7/dist-packages (from
kiwisolver>=1.0.1->matplotlib>=3.0.0->mlxtend) (4.1.1)
Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.7/dist-
packages (from pandas>=0.24.2->mlxtend) (2022.2.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-
packages (from python-dateutil>=2.1->matplotlib>=3.0.0->mlxtend) (1.15.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.7/dist-packages (from scikit-learn>=1.0.2->mlxtend)
(3.1.0)
Installing collected packages: mlxtend
  Attempting uninstall: mlxtend
    Found existing installation: mlxtend 0.14.0
    Uninstalling mlxtend-0.14.0:
      Successfully uninstalled mlxtend-0.14.0
Successfully installed mlxtend-0.21.0

```

## 1.1 Importar módulos de Scikit-learn

```

[ ]: from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, GridSearchCV

from sklearn.neural_network import MLPRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, AdaBoostRegressor,
↳ GradientBoostingRegressor, HistGradientBoostingRegressor
from sklearn.svm import SVR, NuSVR
from sklearn.gaussian_process import GaussianProcessRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn import tree
from sklearn.model_selection import GridSearchCV
import numpy as np
from scipy.stats import loguniform
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
from sklearn.model_selection import RandomizedSearchCV
from mlxtend.evaluate import bias_variance_decomp

```

## 2 Importar Dataset

- **CRIM:** per capita crime rate by town
- **ZN:** proportion of residential land zoned for lots over 25,000 sq.ft.
- **INDUS:** proportion of non-retail business acres per town

- **CHAS**: Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- **NOX**: nitric oxides concentration (parts per 10 million)
- **RM**: average number of rooms per dwelling
- **AGE**: proportion of owner-occupied units built prior to 1940
- **DIS**: weighted distances to five Boston employment centres
- **RAD**: index of accessibility to radial highways
- **TAX**: full-value property-tax rate per \$10,000
- **PTRATIO**: pupil-teacher ratio by town
- **B**:  $1000(B_k - 0.63)^2$  where  $B_k$  is the proportion of blacks by town
- **LSTAT**: % lower status of the population
- **MEDV**: Median value of owner-occupied homes in \$1000's

```
[ ]: url = "https://raw.githubusercontent.com/crisb-7/BostonRealEstate/main/
        ↪bostonRealEstate.csv"
```

```
[ ]: df = pd.read_csv(url)
```

```
[ ]: df.head()
```

```
[ ]:
      CRIM    ZN  INDUS  CHAS    NOX     RM   AGE     DIS  RAD  TAX  PTRATIO  \
0  0.00632  18.0   2.31     0  0.538  6.575  65.2  4.0900    1  296     15.3
1  0.02731   0.0   7.07     0  0.469  6.421  78.9  4.9671    2  242     17.8
2  0.02729   0.0   7.07     0  0.469  7.185  61.1  4.9671    2  242     17.8
3  0.03237   0.0   2.18     0  0.458  6.998  45.8  6.0622    3  222     18.7
4  0.06905   0.0   2.18     0  0.458  7.147  54.2  6.0622    3  222     18.7

      B  LSTAT  MEDV
0  396.90   4.98  24.0
1  396.90   9.14  21.6
2  392.83   4.03  34.7
3  394.63   2.94  33.4
4  396.90   5.33  36.2
```

### 3 Exploración del dataset

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 511 entries, 0 to 510
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   CRIM        511 non-null    float64
1   ZN          511 non-null    float64
2   INDUS       511 non-null    float64
3   CHAS        511 non-null    int64
4   NOX         511 non-null    float64
```

```

5  RM      506 non-null    float64
6  AGE     511 non-null    float64
7  DIS     511 non-null    float64
8  RAD     511 non-null    int64
9  TAX     511 non-null    int64
10 PTRATIO 511 non-null    float64
11 B       511 non-null    float64
12 LSTAT   511 non-null    float64
13 MEDV    511 non-null    float64

```

dtypes: float64(11), int64(3)

memory usage: 56.0 KB

Al ver que se tienen solo 5 registros con valores nulos para RM, se quitan estas filas para tener un conjunto de datos homogéneo.

```
[ ]: df = df.dropna(axis = 0)
df.shape
```

```
[ ]: (506, 14)
```

```
[ ]: df.describe()
```

```
[ ]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.617404	11.289526	11.174842	0.069170	0.555209	6.287589
std	8.600123	23.325350	6.824592	0.253994	0.115611	0.703802
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000
25%	0.082268	0.000000	5.190000	0.000000	0.449000	5.885500
50%	0.266005	0.000000	9.690000	0.000000	0.538000	6.209000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.629750
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000

  

	AGE	DIS	RAD	TAX	PTRATIO	B \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	68.555731	3.775231	9.531621	408.330040	18.498419	356.228379
std	28.161573	2.096147	8.716661	168.382685	2.202078	91.253462
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000
25%	45.025000	2.098500	4.000000	280.250000	17.400000	374.687500
50%	77.500000	3.122200	5.000000	330.000000	19.100000	391.260000
75%	93.975000	5.117675	24.000000	666.000000	20.200000	396.210000
max	100.000000	12.126500	24.000000	711.000000	23.000000	396.900000

  

	LSTAT	MEDV
count	506.000000	506.000000
mean	12.872569	22.711858
std	7.823528	9.520520
min	1.730000	5.000000
25%	6.950000	17.025000

```
50%    11.465000    21.200000
75%    17.107500    25.075000
max     76.000000    67.000000
```

```
[ ]: sns.set(rc={'figure.figsize':(16, 9)})
plt.rcParams["figure.dpi"] = 150
sns.heatmap(df.corr(), annot = True)
plt.show()
```



```
[ ]: # sns.pairplot(df)
# plt.show()
```

## 4 Preprocesamiento de los datos

```
[ ]: scaler = StandardScaler()
```

```
[ ]: x = scaler.fit_transform(df.drop(columns = "MEDV"))
y = df.MEDV
y= y.values
```

## 5 División train-test

```
[ ]: x_train, x_test, y_train, y_test = train_test_split(x, y, train_size = 0.90,
    ↪random_state = 0)
```

## 6 Model Cross-validation

```
[ ]: randomState = 0

Regressors = []

# Regressors.append(MLPRegressor(random_state = randomState, activation =
    ↪"relu", solver = "adam",
#                               hidden_layer_sizes = (100,), alpha = 0.0001,
    ↪learning_rate = "constant",
#                               learning_rate_init = 0.0005, max_iter = 5000))

Regressors.append(MLPRegressor(random_state = randomState, activation = "relu",
    ↪solver = "adam",
                               hidden_layer_sizes = (100,), alpha = 0.0101, learning_rate
    ↪= "adaptive",
                               learning_rate_init = 0.1, max_iter = 1000))

Regressors.append(KNeighborsRegressor(n_neighbors = 2, weights = "uniform", p =
    ↪1))

Regressors.append(DecisionTreeRegressor(random_state=randomState))

Regressors.append(RandomForestRegressor(n_estimators = 250, max_depth = 7,
    ↪random_state=randomState))

Regressors.append(SVR(C = 40.7, epsilon=0.56))

Regressors.append(NuSVR(C = 31.2, nu=0.5))

Regressors.append(AdaBoostRegressor(random_state = randomState))

Regressors.append(GradientBoostingRegressor(random_state = randomState))

Regressors.append(GaussianProcessRegressor(random_state = randomState))

Regressors.append(HistGradientBoostingRegressor(random_state = randomState))

cv_results = []
cv_train_score = []
for regressor in Regressors:
```

```

cv_train=regressor.fit(x_train, y = y_train)
cv_train_score.append(regressor.score(x_train, y_train))
cv_results.append(regressor.score(x_test, y_test))

cv_res = pd.DataFrame({"Algorithm":["NN", "KN","Decision Tree","Random
↳Forest", "SVR", "NuSVR", "AdaBoost", "GradientBoost", "GaussianProcess",
↳"HistGradientBoosting"],
                        "TrainScore":cv_train_score, "TestScore":cv_results})
cv_res.sort_values(by = "TestScore", ascending = False)

```

```

[ ]:

```

	Algorithm	TrainScore	TestScore
7	GradientBoost	0.971574	0.909959
3	Random Forest	0.952958	0.854074
0	NN	0.897917	0.823250
9	HistGradientBoosting	0.958757	0.808177
1	KN	0.929634	0.794065
2	Decision Tree	1.000000	0.778517
4	SVR	0.913308	0.746149
5	NuSVR	0.904299	0.719656
6	AdaBoost	0.862871	0.644497
8	GaussianProcess	1.000000	-0.304075

## 7 Mejora del modelo

Mejoramos el modelo que nos dio más precisión en el test set (GradientBoost) aplicando un tuning de hiperparámetros con ayuda de RandomizedSearchCV

```

[ ]: model = GradientBoostingRegressor()
parameters = {'learning_rate': sp_randFloat(),
              'subsample' : sp_randFloat(),
              'n_estimators' : sp_randInt(50,1000),
              'max_depth' : sp_randInt(2,10)
              }

random = RandomizedSearchCV(estimator=model, param_distributions=parameters, cv
↳= 2, n_jobs=-1)
random.fit(x_train, y_train)

```

```

[ ]: RandomizedSearchCV(cv=2, estimator=GradientBoostingRegressor(), n_jobs=-1,
                        param_distributions={'learning_rate':
<scipy.stats._distn_infrastructure.rv_frozen object at 0x7faf5dc578d0>,
                        'max_depth':
<scipy.stats._distn_infrastructure.rv_frozen object at 0x7faf5dc76b50>,
                        'n_estimators':
<scipy.stats._distn_infrastructure.rv_frozen object at 0x7faf5dc76e90>,
                        'subsample':
<scipy.stats._distn_infrastructure.rv_frozen object at 0x7faf5dc57e10>})

```

```
[ ]: random.best_estimator_  
[ ]: GradientBoostingRegressor(learning_rate=0.03039604699664189, n_estimators=420,  
                               subsample=0.8699095805674916)  
[ ]: random.score(x_test,y_test)  
[ ]: 0.91173883344591  
[ ]: random.score(x_train,y_train)  
[ ]: 0.9801051973743056
```

## 8 Visualización

```
[ ]: pred = {'Real Value':y_test, 'Prediction': random.predict(x_test)}  
      predictions_df = pd.DataFrame(pred)  
  
      plt.plot(predictions_df['Real Value'])  
      plt.plot(predictions_df.Prediction)  
  
      plt.legend(["Real Value", "Prediction"], loc=0)  
      plt.show()
```